

1. OBJECTIVE OF THE PROGRAM

The objective of this code is to perform dimensionality reduction using t-SNE (t-Distributed Stochastic Neighbor Embedding) on the Iris dataset and visualize the dataset in 2D space.

The dataset originally contains 4 features per sample:

1. Sepal length
2. Sepal width
3. Petal length
4. Petal width

Since humans cannot directly visualize 4D data, t-SNE maps the dataset into 2 dimensions, preserving similarity between points as much as possible.

The code also evaluates the embedding quality using:

1. Trustworthiness score
2. Silhouette score

2. KEY TERMS

2.1 Dimensionality Reduction

Definition: Process of reducing the number of features (dimensions) while preserving important structure.

Goal:

Convert high-dimensional dataset $X \in \mathbb{R}^d$ into low-dimensional dataset $Y \in \mathbb{R}^k$, where $k < d$.

2.2 t-SNE (t-Distributed Stochastic Neighbor Embedding)

Definition: A nonlinear dimensionality reduction technique mainly used for visualization. It converts distances between points into probabilities, then tries to preserve those probabilities in low-dimensional space.

t-SNE converts distances into probabilities.

Similarity [High-dimensional] probability: For two points x_i and x_j , t-SNE defines conditional probability:

$$p(j|i) = \frac{\exp\left(-\frac{|x_i - x_j|^2}{2\sigma_i^2}\right)}{\sum_{k \neq i} \exp\left(-\frac{|x_i - x_k|^2}{(2\sigma_i^2)}\right)}$$

Where:

- ◆ $\|x_i - x_j\|^2$ = squared Euclidean distance
- ◆ σ_i = variance controlling neighborhood size

Then symmetric joint probability:

$$p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$$

Similarity [Low-dimensional] probability: In low dimension, similarity is computed using Student t-distribution

$$q_{ij} = \frac{(1 + |y_i - y_j|^2)^{-1}}{\sum_{k \neq l} (1 + |y_k - y_l|^2)^{-1}}$$

This heavy-tailed distribution reduces the crowding problem.

t-SNE minimizes KL divergence (Loss Function): t-SNE minimizes the difference between P and Q using Kullback-Leibler divergence

$$KL(P||Q) = \sum_i \sum_j p_{ij} \log \left(\frac{p_{ij}}{q_{ij}} \right)$$

This ensures that points close in original space remain close in embedding.

2.3 Perplexity (t-SNE Parameter)

Definition: Perplexity controls how many neighbours t-SNE considers important.
Controls the effective number of neighbors.

$$\text{Perplexity}(P_i) = 2^{H(P_i)}$$

Small perplexity is appropriate because the corpus vocabulary is small.

Where entropy is:

$$H(P_i) = - \sum_j p_{j|i} \log_2(p_{j|i})$$

- ◆ Higher perplexity → considers more global structure.
- ◆ Lower perplexity → focuses on local neighborhoods.

2.4 StandardScaler (Standardization)

Definition: Rescales features so that each feature has mean 0 and variance 1.

Formula:

$$z = \frac{x - \mu}{\sigma}$$

Where:

- ◆ μ = mean
- ◆ σ = standard deviation
- ◆ This is required because t-SNE is sensitive to feature scaling.

2.5 Trustworthiness Score

Definition: Measures how well the low-dimensional embedding preserves local neighbors.

Range: 0 to 1 [Higher is better.]

Formula:

$$T(k) = 1 - \frac{2}{nk(2n - 3k - 1)} \sum_{i=1}^n \sum_{j \in U_k(i)} (r(i, j) - k)$$

Where:

- ◆ $U_k(i)$ are points that appear in k-nearest neighbors in embedding but not in original space
- ◆ $r(i,j)$ is rank in original space

2.6 Silhouette Score

Definition: Measures clustering separation.

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

Where:

- ◆ $a(i)$ = average distance to points in same cluster
- ◆ $b(i)$ = average distance to nearest different cluster
- ◆ Range: **-1 to 1** [Higher is better.]

3. CODE EXPLANATION

3.1 Import Libraries

```
import numpy as np
import matplotlib.pyplot as plt
```

`numpy` is imported for numerical operations.

`matplotlib.pyplot` is imported for visualization.

```
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.manifold import TSNE, trustworthiness
from sklearn.metrics import silhouette_score
```

`load_iris` loads the Iris dataset.

`StandardScaler` standardizes features.

`TSNE` runs the t-SNE algorithm.

`trustworthiness` evaluates neighborhood preservation.

`silhouette_score` evaluates clustering quality.

3.2 Load Dataset

```
X, y = load_iris(return_X_y=True)
```

`X` stores features (150 samples \times 4 features).

`y` stores class labels (0, 1, 2).

This is a supervised dataset, but t-SNE itself is unsupervised.

3.3 Standardize Data

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
scaler = StandardScaler() initializes scaling object.
fit_transform() computes mean/std and transforms values.
```

Purpose: Ensures all features contribute equally to distance computations.

3.5 Initialize t-SNE

```
tsne = TSNE(
    n_components=2,
```

```

        perplexity=30,
        learning_rate=200,
        n_iter=1000,
        metric="mahalanobis",
        random_state=42
    )
n_components=2: output will be 2D embedding.
perplexity=30: considers ~30 nearest neighbors.
learning_rate=200: step size for gradient descent updates.
n_iter=1000: number of optimization iterations.
metric="mahalanobis": uses Mahalanobis distance instead of Euclidean.
random_state=42: ensures reproducible results.

```

3.6 Mahalanobis distance formula:

$$d(x, y) = \sqrt{(x - y)^T S^{-1} (x - y)}$$

Where S is covariance matrix.

3.7 Fit and Transform Dataset

```
X_tsne = tsne.fit_transform(X_scaled)
```

Fits t-SNE model on standardized data.
 Returns 2D embedding of shape (150×2) .
 This is the final reduced representation.

3.8 Compute Trustworthiness Score

```
trust = trustworthiness(X_scaled, X_tsne, n_neighbors=5)
print(f"Trustworthiness score: {trust:.4f}")
```

Measures whether nearest neighbors in original space remain neighbors in embedded space.
 $n_{neighbors}=5$ means check neighborhood consistency among top 5 nearest points.

3.9 Compute Silhouette Score

```
sil_score = silhouette_score(X_tsne, y)
print(f"Silhouette score: {sil_score:.4f}")
```

Measures how well the 3 iris classes separate in embedding.
 Uses y labels to check if points of same class are grouped.

3.10 Visualization

```
plt.figure(figsize=(8, 6))
```

Creates a plotting canvas of size 8×6 .

```

scatter = plt.scatter(
    X_tsne[:, 0],
    X_tsne[:, 1],
    c=y,
    cmap="viridis",
    s=60,
    alpha=0.8
)

```

Plots t-SNE output points.

$X_{tsne}[:, 0]$ is x-axis coordinate.

`x_tsne[:,1]` is y-axis coordinate.
`c=y` assigns color by class label.
`cmap="viridis"` sets colormap.
`s=60` controls point size.
`alpha=0.8` makes points slightly transparent.

3.11 Add Colorbar and Labels

```
plt.colorbar(scatter, label="Class Label")  
Adds legend bar mapping colors to class IDs (0,1,2).
```

```
plt.title("t-SNE Visualization (Iris Dataset)")  
plt.xlabel("t-SNE Dimension 1")  
plt.ylabel("t-SNE Dimension 2")  
plt.grid(True)
```

Adds title and axis labels.

`grid(True)` improves readability.

3.12 Display Output

```
plt.show()
```

Renders the scatter plot.

3.13 For 2nd and 3rd part

```
import pandas as pd  
url = "https://raw.githubusercontent.com/pandas-dev/pandas/main/pandas/tests/io/data/csv/tips.csv"  
# Corrected URL to an existing raw CSV file  
data = pd.read_csv(url)  
print(data.head())  
data.info()  
data.describe()  
  
import pandas as pd  
url = "https://raw.githubusercontent.com/pandas-dev/pandas/main/pandas/tests/io/data/csv/iris.csv"  
data = pd.read_csv(url)  
print(data.head())  
data.info()  
data.describe()  
  
.head()
```

Shows sample rows so you can confirm:
dataset is correctly loaded
column names are correct
data types appear meaningful

```
.info()
```

Gives:

- ◆ number of rows and columns
- ◆ datatypes

- ◆ memory usage
- ◆ missing values status

This helps decide preprocessing steps like:

- ◆ encoding categorical features
- ◆ scaling numeric features
- ◆ handling missing values

```
.describe()
```

Gives statistical distribution summary.

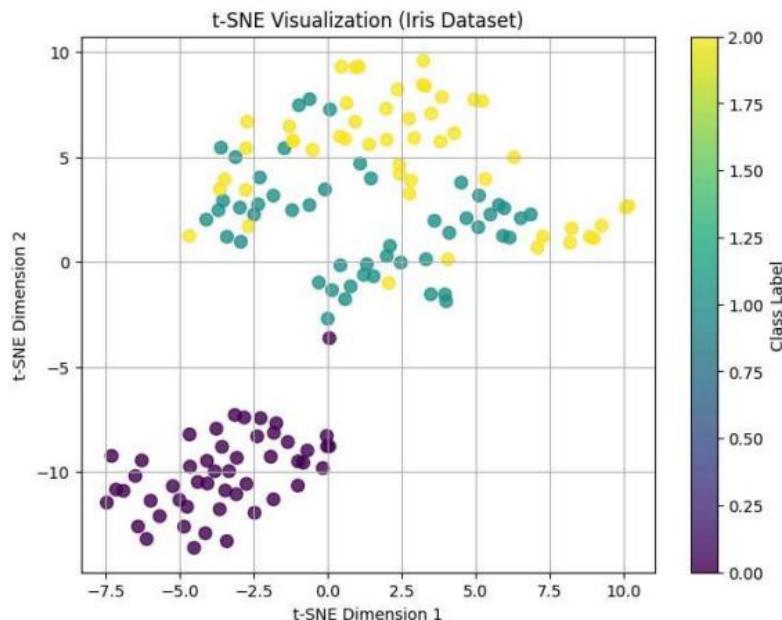
Used for:

- ◆ detecting outliers
- ◆ understanding data spread
- ◆ checking feature ranges

4. Output Interpretation

4.1 Output

```
/usr/local/lib/python3.12/dist-packages/sklearn/manifold/_t_sne.py:1164: FutureWarning: 'n_iter' was renamed to 'max_iter' in ve
warnings.warn(
Trustworthiness score: 0.9361
Silhouette score: 0.3351
```



The output includes:

Trustworthiness score: 0.9361

Silhouette score: 0.3351

Scatter plot titled: *t-SNE Visualization (Iris Dataset)*

Trustworthiness Score Interpretation (0.9361)

- ◆ This value is close to 1.
- ◆ It indicates that local neighborhood structure is strongly preserved.
- ◆ Meaning: nearest neighbors in original 4D space remain neighbors in 2D.

This is considered a strong embedding.

Silhouette Score Interpretation (0.3351)

- ◆ Value is moderately positive.
- ◆ This means clusters exist but are not perfectly separated.
- ◆ For Iris dataset, one class (Setosa) is usually well separated, while the other two overlap. So, this score is logically consistent.

Scatter Plot Explanation

- ◆ The purple cluster (label 0) forms a clear separated group at the bottom-left region. This corresponds to Iris Setosa, which is naturally separable.
- ◆ The green and yellow points (labels 1 and 2) overlap heavily in the upper region. These represent Iris Versicolor and Iris Virginica, which share similar feature patterns.
- ◆ The axes represent t-SNE coordinates, not original features.
- ◆ Their numeric scale has no direct physical meaning.
- ◆ The main interpretation is based on relative closeness, not absolute axis values.

t-SNE pipeline:

preprocessing → reduction → evaluation → visualization.

4.2 Output

```
total_bill    tip     sex smoker   day    time   size
0      16.99  1.01  Female    No Sun Dinner     2
1      10.34  1.66   Male    No Sun Dinner     3
2      21.01  3.50   Male    No Sun Dinner     3
3      23.68  3.31   Male    No Sun Dinner     2
4      24.59  3.61  Female    No Sun Dinner     4
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244 entries, 0 to 243
Data columns (total 7 columns):
 #   Column   Non-Null Count  Dtype  
 --- 
 0   total_bill    244 non-null   float64
 1   tip          244 non-null   float64
 2   sex          244 non-null   object  
 3   smoker        244 non-null   object  
 4   day           244 non-null   object  
 5   time          244 non-null   object  
 6   size          244 non-null   int64   
dtypes: float64(2), int64(1), object(4)
memory usage: 13.5+ KB
      total_bill      tip      size
count  244.000000  244.000000  244.000000
mean   19.785943  2.998279  2.569672
std    8.902412  1.383638  0.951100
min    3.070000  1.000000  1.000000
25%   13.347500  2.000000  2.000000
50%   17.795000  2.900000  2.000000
75%   24.127500  3.562500  3.000000
max    50.810000 10.000000  6.000000
```

This code does not create graphs, but it produces tabular outputs

total_bill

```
count = 244
mean ≈ 19.7859
std ≈ 8.9024
min = 3.07
max = 50.81
```

Meaning: Bills range from 3.07 to 50.81.

tip

```
mean ≈ 2.9983
max = 10.00
```

Meaning: Average tip is ~3.

size

mean \approx 2.5697

max = 6

Meaning: Average group size is ~2.6 people.

4.3 Output

```
SepalLength SepalWidth PetalLength PetalWidth      Name
0          5.1        3.5       1.4       0.2 Iris-setosa
1          4.9        3.0       1.4       0.2 Iris-setosa
2          4.7        3.2       1.3       0.2 Iris-setosa
3          4.6        3.1       1.5       0.2 Iris-setosa
4          5.0        3.6       1.4       0.2 Iris-setosa
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   SepalLength  150 non-null    float64 
 1   SepalWidth   150 non-null    float64 
 2   PetalLength  150 non-null    float64 
 3   PetalWidth   150 non-null    float64 
 4   Name         150 non-null    object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
SepalLength  SepalWidth  PetalLength  PetalWidth
count      150.000000  150.000000  150.000000  150.000000
mean       5.843333   3.054000   3.758667   1.198667
std        0.828066   0.433594   1.764420   0.763161
min        4.300000   2.000000   1.000000   0.100000
25%        5.100000   2.800000   1.600000   0.300000
50%        5.800000   3.000000   4.350000   1.300000
75%        6.400000   3.300000   5.100000   1.800000
max        7.900000   4.400000   6.900000   2.500000
```

This code does not create graphs, but it produces tabular outputs

SepalLength

mean \approx 5.8433

min = 4.3

max = 7.9

SepalWidth

mean \approx 3.0540

min = 2.0

max = 4.4

PetalLength

mean \approx 3.7587

min = 1.0

max = 6.9

PetalWidth

mean \approx 1.1987

min = 0.1

max = 2.5

These values represent physical measurements of iris flower parts.

5. Analytical Recapitulation

This code integrates two fully developed workflows: (1) loading the dataset and performing exploratory data analysis (EDA) with Pandas on both tips.csv and iris.csv, and (2) applying dimensionality reduction and visualizing the Iris dataset using t-SNE. To start, the program imports essential Python libraries, including Pandas for data manipulation, NumPy for handling numerical operations, Matplotlib for graphical representation, and various scikit-learn modules for dataset loading, feature scaling, t-SNE application, and assessment metrics. The initial workflow retrieves the tips.csv dataset from a raw GitHub URL via `pd.read_csv(url)` and saves it into a Pandas DataFrame. The code subsequently displays the first few rows with `.head()` to confirm successful loading and examine sample values. Following this, `.info()` provides an overview of the dataset's structure, revealing a total of 244 rows and 7 columns, along with data types, non-null counts, and memory information. The dataset comprises numeric columns like `total_bill`, `tip`, and `size`, alongside categorical columns such as `sex`, `smoker`, `day`, and `time`. Lastly, `.describe()` generates summary statistics for the numeric columns, including count, mean, standard deviation, minimum, maximum, and quartiles. These statistics are based on mathematical formulas such as mean $\mu = \frac{1}{n} \sum x_i$ and standard deviation

$$\sigma = \sqrt{\frac{1}{n-1} \sum (x_i - \mu)^2}$$

helping identify data distribution and possible outliers. The second workflow involves loading the Iris dataset, which can be done through either scikit-learn's `load_iris(return_X_y=True)` function or by importing the `iris.csv` file from GitHub's raw URL into a DataFrame. This process generates a feature matrix `(X)` that comprises measurements of the flowers and a label vector `(y)` that holds the class labels. The code utilizes `.head()`, `.info()`, and `.describe()` to analyze the dataset's structure, verifying that there are 150 entries and 5 columns in the CSV version. Among these, four columns represent numerical measurements of the flowers (`SepalLength`, `SepalWidth`, `PetalLength`, `PetalWidth`), while one column indicates the categorical label `Name`. After inspection, the t-SNE pipeline begins by standardizing the Iris feature matrix using `StandardScaler`, applying normalization $z = (x - \mu)/\sigma$ so that each feature has mean 0 and standard deviation 1, which is necessary because t-SNE is highly sensitive to feature scaling. The t-SNE algorithm is initialized with various parameters: `n_components=2` for a 2D embedding, `perplexity=30` to manage the size of neighborhoods, `learning_rate=200` to dictate the size of the gradient descent steps, and `n_iter=1000` for the number of optimization iterations. It uses Mahalanobis distance as the metric for similarity measurement and sets `random_state=42` to ensure reproducibility. The t-SNE process transforms high-dimensional pairwise distances into probability distributions `(p_{ij})`, calculates low-dimensional similarity probabilities `(q_{ij})` using a Student t-distribution to alleviate the crowding issue, and aims to minimize the difference between `(P)` and `(Q)` through the KL divergence `(KL(P || Q))` via gradient descent optimization. Once the final 2D embedding is produced, the code assesses the quality of the dimensionality reduction by calculating a trustworthiness score of 0.9361, which indicates a strong retention of local neighborhood relationships from the original space, along with a silhouette score of 0.3351, suggesting moderate separation among the three iris classes. Finally, the reduced 2D points are visualized using a scatter plot where each point represents a flower sample and colors represent class labels. The visualization shows one clearly separated cluster (`Setosa`) and overlapping clusters for the other two classes (`Versicolor` and `Virginica`), which aligns with known Iris dataset behavior. Overall, the combined code demonstrates a complete workflow starting from online dataset loading and statistical inspection, followed by preprocessing, dimensionality reduction using t-SNE, quantitative embedding evaluation, and final 2D visualization.